

# EOLE: Combining Static and Dynamic Scheduling Through Value Prediction to Reduce Complexity and Increase Performance

ARTHUR PERAIS and ANDRÉ SEZNEC, INRIA

Recent work in the field of value prediction (VP) has shown that given an efficient confidence estimation mechanism, prediction validation could be removed from the out-of-order engine and delayed until commit time. As a result, a simple recovery mechanism—pipeline squashing—can be used, whereas the out-of-order engine remains mostly unmodified.

Yet, VP and validation at commit time require additional ports on the physical register file, potentially rendering the overall number of ports unbearable. Fortunately, VP also implies that many single-cycle ALU instructions have their operands predicted in the front-end and can be executed in-place, in-order. Similarly, the execution of single-cycle instructions whose result has been predicted can be delayed until commit time since predictions are validated at commit time.

Consequently, a significant number of instructions—10% to 70% in our experiments—can bypass the out-of-order engine, allowing for a reduction of the issue width. This reduction paves the way for a truly practical implementation of VP. Furthermore, since VP in itself usually increases performance, our resulting {Early—Out-of-Order—Late} Execution architecture, EOLE, is often more efficient than a baseline VP-augmented 6-issue superscalar while having a significantly narrower 4-issue out-of-order engine.

CCS Concepts: • **Computer systems organization** → **Superscalar architectures**; *Complex instruction set computing*; *Pipeline computing*;

Additional Key Words and Phrases: Value prediction, speculative execution, out-of-order execution, EOLE, VTAGE

## ACM Reference Format:

Arthur Perais and André Seznec. 2016. EOLE: Combining static and dynamic scheduling through value prediction to reduce complexity and increase performance. *ACM Trans. Comput. Syst.* 34, 2, Article 4 (April 2016), 33 pages.

DOI: <http://dx.doi.org/10.1145/2870632>

## 1. INTRODUCTION AND MOTIVATIONS

Even in the multicore era, the need for higher single thread performance is driving the definition of new high-performance cores. Although the usual superscalar design does not scale, increasing the ability of the processor to extract instruction-level parallelism (ILP) by increasing the window size and the issue width has generally been the favored way to enhance sequential performance. For instance, consider the recently introduced Intel Haswell microarchitecture that has 33% more issue capacity than Intel Nehalem.<sup>1</sup> To accommodate this increase, both the reorder buffer (ROB) and scheduler size were

<sup>1</sup>State of the art in 2009.

This work was partially supported by the European Research Council Advanced Grant DAL 267175.

Authors' addresses: A. Perais, Bureau 311, Batiment E, Campus de Beaulieu, 263 Avenue Général Leclerc, 35042 Rennes, France; email: [arthur.perais@inria.fr](mailto:arthur.perais@inria.fr); A. Seznec, Bureau 321, Batiment E, Campus de Beaulieu, 263 Avenue Général Leclerc, 35042 Rennes, France; email: [Andre.Seznec@inria.fr](mailto:Andre.Seznec@inria.fr).

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

© 2016 Copyright held by the owner/author(s).

0734-2071/2016/04-ART4 \$15.00

DOI: <http://dx.doi.org/10.1145/2870632>

substantially increased.<sup>2</sup> On top of this, modern schedulers must support complex mechanisms such as *speculative scheduling* to enable back-to-back execution and thus *replay* or *selective replay* to efficiently recover from schedule mispredictions [Kim and Lipasti 2004].

In addition, the issue width impacts other structures: the physical register file (PRF) must provision more read/write ports as the width grows, whereas the number of physical registers must also increase to accommodate the ROB size. Because of this, both latency and power consumption increase and using a monolithic register file quickly becomes complexity ineffective. Similarly, a wide-issue processor should provide enough functional units to limit resource contention. Yet, the complexity of the bypass network grows quadratically with the number of functional units and quickly becomes critical regarding cycle time [Palacharla et al. 1997]. In other words, the out-of-order engine impact on power consumption and cycle time is ever increasing [Ernst and Austin 2002].

In this article, we propose a modified superscalar design—the {Early—Out-of-Order—Late} Execution microarchitecture, EOLE.<sup>3</sup> It is built on top of a value prediction (VP) pipeline. VP allows dependents to issue earlier than previously possible by using predicted operands, artificially increasing ILP. Yet predictions must be verified to ensure correctness. Fortunately, Perais and Seznec observed that the cost of validating the predicted results (and recovering from mispredictions) at retirement can be absorbed provided an enhanced confidence estimation mechanism that yields a very high prediction accuracy [Perais and Seznec 2014b]. In other words, VP does not need to intervene in the execution engine, save for the PRF.

With EOLE, we leverage this observation to further reduce both the complexity of the out-of-order execution engine and the number of ports required on the PRF when VP is implemented. We achieve this reduction without significantly impacting overall performance. Our contribution is therefore twofold. First, EOLE paves the way to truly practical implementations of VP. Second, it reduces complexity in arguably the most complex and power-hungry part of a modern superscalar core.

EOLE relies on the fact that when using VP, a significant number of single-cycle instructions have their operands ready in the front-end thanks to the value predictor. As such, we introduce *early execution* (EE) to execute single-cycle ALU instructions in-order in parallel with Rename by using predicted and/or immediate operands. Early-executed instructions are *not* sent to the out-of-order scheduler. Moreover, delaying VP validation until commit time removes the need for *selective replay* and enforces a complete pipeline squash on a value misprediction. This guarantees that the operands of committed early-executed instructions were the correct operands. EE requires simple hardware and reduces pressure on the out-of-order instruction window.

Similarly, since predicted results can be validated outside the out-of-order engine at commit time [Perais and Seznec 2014b], we can offload the execution of predicted single-cycle ALU instructions to some dedicated in-order *late execution* (LE) precommit stage, where no *Select & Wakeup* has to take place. This does not hurt performance, as instructions dependent on predicted instructions will simply use the predicted results rather than wait in the out-of-order scheduler. Similarly, the resolution of high-confidence branches can be offloaded to the LE stage since they are very rarely mispredicted.

Overall, a total of 10% to 70% of the retired instructions can be offloaded from the out-of-order engine. As a result, EOLE benefits from both the aggressiveness of modern

<sup>2</sup>From 128 and 36 entries to 192 and 60 entries, respectively.

<sup>3</sup>*Eole* is the French name of *Aeolus*, the ruler of the winds in Greek mythology.

out-of-order designs and the higher energy efficiency of more conservative in-order designs.

We evaluate EOLE against a baseline out-of-order model featuring VP and show that it achieves similar levels of performance having only 66% of the baseline issue capacity and a significantly less complex PRF. This is especially interesting because it provides architects extra design headroom in the out-of-order engine to implement new architectural features.

The remainder of this article is organized as follows. Section 2 discusses related work and provides some background on VP. Section 3 details the EOLE microarchitecture, which implements both EE and LE by leveraging VP. Section 4 describes our simulation framework, and Section 5 presents experimental results. Section 6 focuses on the qualitative gains in complexity and power consumption permitted by EOLE. Finally, Section 7 provides concluding remarks and directions for future research.

## 2. RELATED WORK

Many propositions aim at reducing complexity in modern superscalar designs. In particular, it has been shown that most of the complexity and power consumption reside in the out-of-order engine, including the PRF [Wallace and Bagherzadeh 1996], scheduler, and bypass network [Palacharla et al. 1997]. As such, previous studies focused either on devising new pipeline organizations or reducing the complexity of existing structures.

### 2.1. Alternative Pipeline Organizations

Farkas et al. [1997] propose the *multicluster* architecture in which execution is distributed among several execution clusters, each of them having its own register file. Since each cluster is simpler, cycle time can be decreased even though some inefficiencies are introduced due to intercluster dependencies. In Farkas et al. [1997], intercluster data dependencies are handled by dispatching the same instruction to several clusters and having all instances but one serve as intercluster data-transfer instructions while a single instance actually computes the result. To enforce correctness, this instance is data dependent on all others. The Alpha 21264 [Kessler et al. 1998] is an example of real-world clustered architecture and shares many traits with the multicluster architecture.

Palacharla et al. [1997] introduce a *dependence-based* microarchitecture where the centralized instruction window is replaced by several parallel FIFOs. FIFOs are filled with independent instructions (i.e., all instructions in a single FIFO are dependent) to allow ILP to be extracted: since instructions at the head of the FIFOs are generally independent, they can issue in parallel. This greatly reduces complexity since only the head of each FIFO has to be scanned by the *Select* logic. They also study a *clustered dependence-based* architecture to reduce the amount of bypass and window logic by using clustering. In other words, a few FIFOs are grouped together and assigned their own copy of the register file and bypass network, mimicking a wide issue window by having several smaller ones. Intercluster bypassing is naturally slower than intracluster bypassing.

Tseng and Patt [2008] propose the *Braid* architecture, which shares many similarities with the clustered dependence-based architecture with three major differences. First, instruction partitioning—steering—is done at compile time via dataflow-graph coloring. Second, each FIFO is a narrow-issue (dual-issue in Tseng and Patt [2008]) in-order cluster called a *Braid Execution Unit* with its own local register file, execution units, and bypass network. Third, a global register file handles interunit dependencies instead of an intercluster bypass network. As such, they obtain performance on par

with an accordingly sized out-of-order machine without using complicated scheduling logic, register file, and bypass logic.

Austin [1999] proposes *Dynamic Implementation VALidation* (DIVA) to check instruction results just before commit time, allowing the core to be faulty. In other words, DIVA can greatly reduce the complexity of *functionally validating* the core since only the checker has to be functionally correct. The core complexity itself can remain very high, however. An interesting observation is that the latency of the checker has very limited impact on performance. This hints that adding pipeline stages between Write-back and Commit (such as VP validation) should not impact performance much.

Fahs et al. [2005] study *continuous optimization*, where common compile-time optimizations are applied dynamically in the rename stage. This allows early execution of some instructions in the front-end instead of the out-of-order engine. Similarly, Petric et al. [2005] propose RENO, which also dynamically applies optimizations at rename time.

## 2.2. Decreasing the Complexity of Implemented Mechanisms

Instead of studying new organizations of the pipeline, Kim and Lipasti [2003] present the *half-price architecture*. They argue that many instructions are single-operand and that both operands of dual-operands instructions rarely become ready at the same time. Thus, the load capacitance on the tag broadcast bus can be greatly reduced by sequentially waking up operands. In particular, the left operand is woken up as usual, whereas the right one is woken up one cycle later by inserting a latch on the broadcast bus. This scheme relies on an operand criticality predictor to place the critical tag in the left operand field.

Similarly, Ernst and Austin [2002] propose *tag elimination* to limit the number of comparators used for Wakeup. In particular, only the tag of the last arriving operand (predicted as such) is put in the scheduler. This allows the use of one comparator per entry instead of several (one per operand). However, this also implies that the instruction must be replayed on a wrong last operand prediction.

Regarding the PRF, Kim and Lipasti [2003] also observe that many issuing instructions do not need to read both their operands in the PRF, as one or both will be available on the bypass network. Thus, provisioning two read ports per issue slot is generally overprovisioning. In particular, for the small proportion of instructions that indeed need to read two operands in the PRF, they propose to do two consecutive accesses using a single port. Reducing the number of ports drastically reduces the complexity of the register file, as ports are much more expensive than registers.

Last, Lukefahr et al. [2012] propose to implement two back-ends—in-order and out-of-order—in a single core and to dynamically dispatch instructions to the most adapted one. In most cases, this saves power at a slight cost in performance. In a sense, EOLE has similarities with such a design, as instructions can be executed in different locations. However, no decision has to be made regarding the location where an instruction will be executed, because this depends only on the instruction type and status (e.g., predicted or not).

Note that our proposal is orthogonal to all of these contributions, as it only impacts the number of instructions that enters the out-of-order execution engine.

*Value prediction.* EOLE builds upon the broad spectrum of research on VP independently initiated by Gabbay and Mendelson [1998] and Lipasti and Shen [1996].

Sazeides and Smith [1997] refine the taxonomy of VP by categorizing predictors. They define two classes of value predictors: *computational* and *context based*. Computational predictors generate a prediction by applying a function to the value(s) produced by the previous instance(s) of the instruction. For example, the Stride predictor [Mendelson

and Gabbay 1997] and the 2-Delta Stride predictor [Eickemeyer and Vassiliadis 1993] use the addition of a constant (stride).

On the other hand, context-based predictors rely on patterns in the value history of a given static instruction to generate predictions, such as the finite context method (FCM) predictors [Sazeides and Smith 1997].

Most of the initial studies on VP either assume that recovering from a value misprediction induces almost no penalty [Lipasti and Shen 1996; Lipasti et al. 1996; Zhou et al. 2003] or simply focus on accuracy and coverage rather than speedup [Goeman et al. 2001; Nakra et al. 1999; Rychlik et al. 1998; Sazeides and Smith 1997; Thomas and Franklin 2001; Wang and Franklin 1997]. The latter studies were essentially ignoring the performance loss associated with misprediction recovery.

In a recent study, Perais and Seznec [2014b] show that all value predictors are amenable to very high accuracy at a reasonable cost in both prediction coverage and hardware storage. This allows the delay of prediction validation until commit time, removing the burden of implementing a complex replay mechanism that is tightly coupled to the out-of-order engine [Kim and Lipasti 2004]. As such, the out-of-order engine remains mostly untouched by VP. This proposition is crucial, as VP was usually considered very hard to implement in part due to the need for a very fast recovery mechanism.

In the same paper [Perais and Seznec 2014b], the VTAGE context-based predictor is introduced. As the ITTAGE indirect branch predictor [Seznec and Michaud 2006], VTAGE uses the global branch history to select predictions, meaning that it does not require the previous value to predict the current one. This is a strong advantage because conventional value predictors usually need to track inflight predictions, as they require the last value to predict.

Finally, Perais and Seznec [2015] propose a tightly coupled hybrid of a VTAGE predictor and a Stride-based predictor. Similarly to the *differential* FCM (D-FCM) predictor of Goeman et al. [2001], the *differential* VTAGE predictor (D-VTAGE) implements a last value table (LVT) containing the last outcome of each static instruction, to which a stride is added to compute the prediction. The stride is retrieved using the TAGE/VTAGE indexing scheme. D-VTAGE has been shown to outperform a similarly sized D-FCM predictor as well as a more simple hybrid selecting between VTAGE and 2-Delta Stride based on confidence. A practical implementation of the speculative window required to track inflight last outcomes is also provided in Perais and Seznec [2015].

### 3. EOLE

#### 3.1. Enabling EOLE Through Value Prediction

As described previously, EOLE consists of a set of simple ALUs in the in-order front-end to early-execute instructions in parallel with Rename and a second set in the in-order back-end to late execute instructions just before they are committed.

Although EOLE is heavily dependent on VP, they are in fact complementary features. Indeed, the former needs a value predictor to predict operands for EE and provide temporal slack for LE, whereas VP needs EOLE to reduce PRF complexity and thus become truly practical.

Moreover, to be implemented, EOLE requires prediction validation to be done at Commit since validating at Execute mechanically forbids LE. In addition, using *selective replay* to recover from a value misprediction nullifies the interest of both EE and LE, as all instructions must flow through the out-of-order scheduler in case they need to be replayed [Kim and Lipasti 2004]. Hence, *squashing* must be used to recover from a misprediction so that early/late-executed instructions can safely bypass the scheduler.



Fortunately, Perais and Seznec [2014b] have proposed a confidence estimation mechanism greatly limiting the number of value mispredictions, forward probabilistic counters (FPC). With FPC, the cost of a single misprediction can be high, as mispredicting is very rare. Thus, validation can be done late—at commit time—and squashing can be used as the recovery mechanism. This enables the implementation of both EE and LE, hence EOLE.

By eliminating the need to dispatch and execute many instructions in the out-of-order engine, EOLE substantially reduces the pressure put on complex and power-hungry structures. Thus, those structures may be scaled down, yielding a less complex architecture whose performance is on par with a more aggressive design. Moreover, doing so is orthogonal to previously proposed mechanisms such as *clustering* [Farkas et al. 1997; Kessler et al. 1998; Palacharla et al. 1997; Seznec et al. 2002] and does not *require* a centralized instruction window, even though this is the model that we use in this article. Figure 1 depicts the EOLE architecture, implementing EE (red), LE (green), and VP (orange). In the following paragraphs, we detail the two additional blocks required to implement EOLE and their interactions with the rest of the pipeline.

### 3.2. Early Execution Hardware

The core idea of EE is to position one or more ALU stages in the front-end in which instructions with available operands will be executed. For complexity concerns, however, it seems necessary to limit EE to single-cycle ALU instructions. Indeed, implementing complex functional units in the front-end to execute multicycle instructions does not appear as a worthy trade-off. In particular, memory instructions are not early executed. EE is done in-order, and hence it does not require renamed registers and can take place in parallel with Rename. For instance, Figure 2 depicts the EE block adapted to a 2-wide rename stage.


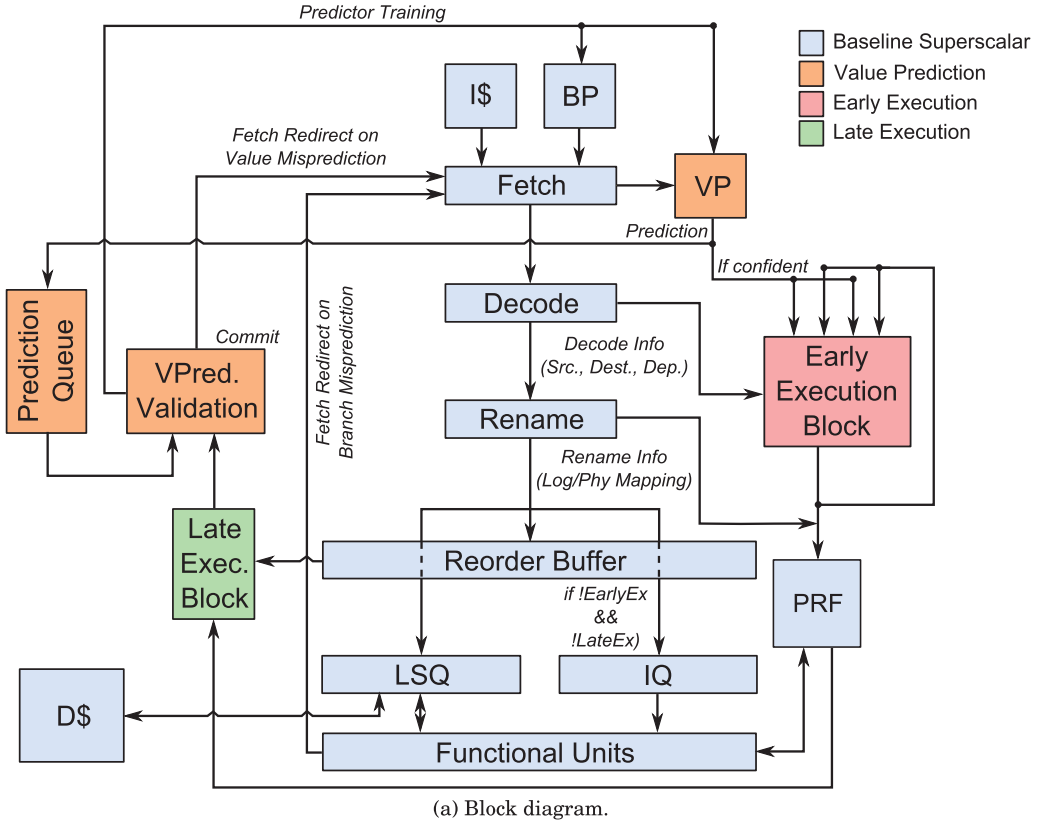
Renaming is often pipelined over several cycles. Consequently, we can use several ALU stages and simply insert pipeline registers between each stage. The actual execution of an instruction can then happen in any of the ALU stages, depending on the readiness of its operands coming from *Decode* (i.e., immediate), the local<sup>4</sup> bypass network (i.e., from instructions early executed in the previous cycle), or the value predictor. Operands are *never* read from the PRF.

In a nutshell, all eligible instructions flow through the ALU stages, propagating their results in each bypass network accordingly once they have executed. Finally, after the last stage, results and predictions are written into the PRF.

An interesting design concern lies with the number of stages required to capture a reasonable proportion of instructions. We actually found that using more than a single stage was highly inefficient, as illustrated in Figure 3. This figure shows the proportion of committed instructions eligible for EE for a baseline 8-wide rename, 6-issue model (see Table I in Section 4), using the D-VTAGE value predictor (later described in Table II, Section 4). As a result, in further experiments, we consider a 1-deep EE block only.

To summarize, EE only requires a single new block, which is shown in red in Figure 1. The mechanism that we propose does not require any storage area for temporaries, as all values are living inside the pipeline registers or the bypass network(s). Finally, since we execute in-order, each instruction is mapped to a single ALU and scheduling is straightforward (as long as there are as many ALUs as the rename width).

<sup>4</sup>For complexity concerns, we consider that bypass does not span several ALU stages. Consequently, if an instruction depends on a result computed by an instruction located two rename groups ahead, it will not be early executed.

Fig. 1. The EOLE  $\mu$ -architecture.

### 3.3. Late Execution Hardware

LE targets instructions whose result has been predicted. Instructions eligible for prediction are  $\mu$ -ops producing a 64-bit or less result that can be read by a subsequent  $\mu$ -op, including nonarchitectural temporary registers, as defined by the ISA implementation. Unless mentioned otherwise, vector and scalar SIMD instructions are not predicted.

LE intervenes just before prediction validation time—that is, out of the execution engine. As for EE, we limit ourselves to single-cycle ALU instructions to minimize complexity—that is, predicted loads are executed in the out-of-order engine but validated at Commit.

Interestingly, Seznez [2011] showed that conditional branch predictions flowing from TAGE can be categorized such that very high confidence predictions are known. Since high-confidence branches exhibit a misprediction rate generally lower than

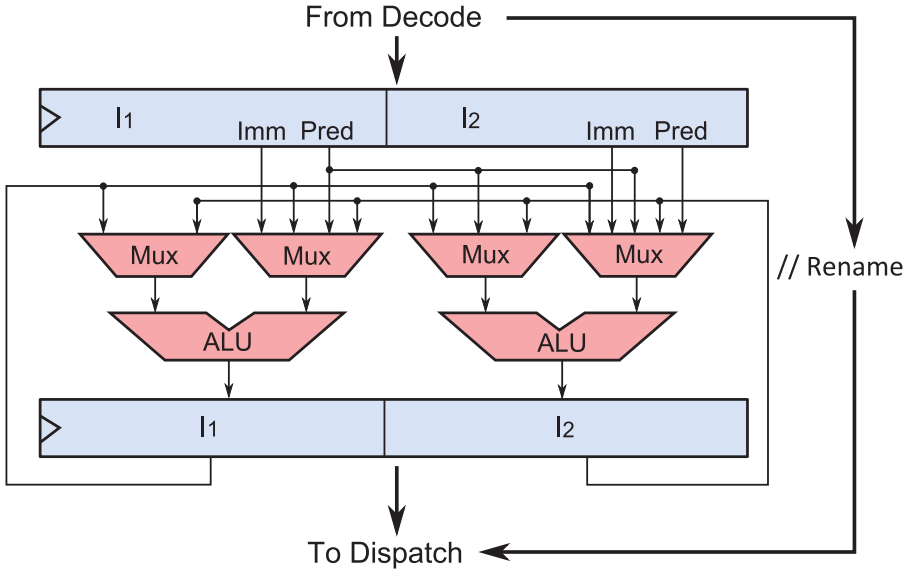


Fig. 2. EE block. The logic controlling the ALUs and muxes is not shown for clarity.




Fig. 3. Proportion of committed instructions that can be early executed, using one or two ALU stages and a D-VTAGE hybrid predictor (later described in Section 4).

0.5%, resolving them in the LE block will have a marginal impact on overall performance. Thus, we consider both single-cycle predicted ALU instructions and very high confidence branches<sup>5</sup> for LE. In this study, we did not try to set confidence on the other branches (indirect jumps, returns). Yet, provided a similar high confidence estimator for these categories of branches, one could postpone the resolution of high confidence ones until the LE stage.

Furthermore, note that predicted instructions can also be early executed. In that event, they only need to be validated in case another early-executed instruction from the same rename group used the prediction as an operand. In other words, instructions are executed in a single location.

In any case, LE further reduces pressure on the out-of-order engine in terms of instructions dispatched to the scheduler. As such, it also removes the need for value predicting only critical instructions [Fields et al. 2001; Rychlik et al. 1998; Tune et al. 2002], as minimizing the number of instructions flowing through the out-of-order engine requires maximizing the number of predicted instructions. Hence, predictions

<sup>5</sup>Predictions whose confidence counter is saturated [Seznec 2011].






Fig. 4. Proportion of committed instructions that can be late executed using a D-VTAGE (see Section 4) hybrid predictor. Late-executable instructions that can also be early executed are not counted, as instructions are executed at most once.

considered as useless from a performance standpoint become useful in EOLE. Figure 4 shows the proportion of committed single-cycle ALU instructions that can be late executed using a baseline 6-issue processor with a D-VTAGE predictor (respectively described in Tables I and II in Section 4).

LE needs to implement *commit width* ALUs and the associated read ports in the PRF. If an instruction  $I_1$  to be late executed depends on the result of instruction  $I_0$  of the same commit group that will also be late executed, then it does not need to wait, as it can use the predicted result of  $I_0$ . In other words, all nonexecuted instructions reaching the LE stage have all of their operands ready in the PRF, as in DIVA [Austin 1999]. Due to the need to validate predictions (including reading results to train the value predictor) and to late execute some instructions, at least one extra pipeline stage after Writeback is likely to be required in EOLE. In the remainder of this article, we refer to this stage as the LE/validation and training (LE/VT) stage.

Overall, the hardware needed for LE is fairly simple, as suggested by the high-level view of a 2-wide LE block shown in Figure 5. It does not even require a bypass network. In further experiments, we consider that LE and prediction validation can be done in the same cycle, before the commit stage. EOLE is therefore only one cycle longer than the baseline superscalar to which it is compared. Although this may be optimistic due to the need to read from the PRF, this only impacts the value misprediction penalty, the pipeline fill delay, and ROB occupancy. In particular, since low-confidence branches are resolved in the same cycle as for the baseline, the average branch misprediction penalty will remain very similar. Last, as a first step, we also consider that enough ALUs are implemented (i.e., as many as the commit width). As a second step, we shall consider reduced-width LE.

### 3.4. Potential Out-of-Order Engine Offload

**3.4.1. Baseline.** We obtain the ratio of retired instructions that can be offloaded from the out-of-order engine for each benchmark by summing the columns in Figures 3 and 4 (both sets are disjoint, as we only count late-executable instructions that cannot also be early executed) and adding *load immediate* instructions. Those numbers are reported in Figure 6, where we distinguish four types of instructions:

- (1) Early executed because it is a load immediate (i.e., its operand is necessarily ready at Rename).<sup>6</sup> In a regular VP pipeline, these instructions can also bypass the out-of-order engine, as the immediate can be considered as an always correct VP. This includes loading an immediate to a floating-point (FP) register.
- (2) Early executed because it is a ready single-cycle ALU instruction.

<sup>6</sup>Except for load immediate to a partial register.

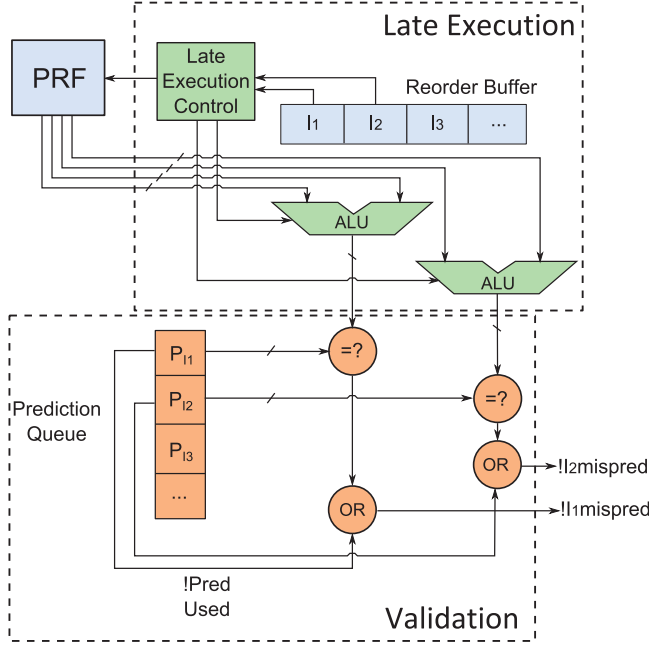


Fig. 5. LE block for a 2-wide processor. The top part can late execute two instructions, whereas the bottom part validates two results against their respective predictions. Buses are *general-purpose register (GPR)*-width-bit wide.




Fig. 6. Proportion of dynamic instructions that can bypass the execution engine by being a high-confidence branch, a value-predicted ALU instruction, an ALU instruction with its operands ready in the front-end, or a load immediate instruction.

- (3) Late executed because it is a high-confidence branch, as predicted by the TAGE branch predictor.
- (4) Late executed because it is a value-predicted single-cycle ALU instruction.

We observe that the ratio is very dependent on the application, ranging from less than 10% for *equake*, *milc*, and *lbm* to more than 50% for *swim*, *mgrid*, *applu*, *perlbench*, *leslie*, *GemsFDTD*, and *xalancbmk* and up to 60% for *namd* and 70% for *libquantum*. It generally represents a significant part of the retired instructions in most cases (around 35% on average).

**3.4.2. Limit Study.** We also consider numbers when we do not specify any constraint regarding the type of instructions that can be early or late executed. In particular, all load and complex arithmetic (e.g., multiply, divide) instructions are considered for bypassing the execution engine. We also consider FP operations and packed vector




Fig. 7. Proportion of dynamic instructions that can bypass the execution engine, depending on what kind of instructions are value predicted and what kind of instructions are allowed to be early/late executed.

operations (the shortcomings of actually predicting those types of instructions are discussed in the next section).

Figure 7 shows the overall potential either without FP/SIMD prediction (but still allowing loads and complex integer instructions to be early or late executed second set of bars) and with FP/SIMD prediction (third set of bars). The first set of bars recalls the numbers of Figure 6.

The second set of bars (no FP/SIMD prediction but all instructions are eligible for EE/LE) shows that in most cases, allowing all instructions to be early/late executed significantly increases the proportion of instructions that can bypass the execution engine. However, this increase is generally due to an increase in late-executable instructions (from 14.3% to 22.3% on average) rather than early-executable ones (from 14.1% to 14.5% on average). Moreover, in most cases, the bulk of the additional candidates for LE consists of value-predicted loads. Indeed, on average, loads represent 26.2% of the predicted instructions versus 1.56% for other types of instructions (multiplications, divisions, and conversions from *float* to *int*) that are now allowed to be early/late executed. As a result, allowing loads to be late executed would be of great help to further reduce pressure on the out-of-order engine.

The third set of bars (FP/SIMD prediction, all instructions are eligible for EE/LE) suggests that although the value predictor is able to predict a significant amount of instructions writing to *xmm* registers for a few benchmarks (e.g., *equake*, *milc*, and *soplex*), the average increase in early/late-executable instructions is much lower than the one we observed by allowing loads and complex integer instructions to be early/late executed.

## 4. EVALUATION METHODOLOGY

### 4.1. Simulator

We use the x86\_64 ISA to validate EOLE, even though EOLE can be adapted to any general-purpose ISA without any more restrictions than the ones already existing for VP. We use a modified<sup>7</sup> version of the *gem5* cycle-level simulator [Binkert et al. 2011].

<sup>7</sup>Our modifications mostly lie with the ISA implementation. In particular, we implemented branches with a single  $\mu$ -op instead of three and removed some false dependencies existing between instructions due to the way flags are renamed/written.

Table I. Simulator Configuration Overview

Front-End	L1I 8-way 32KB, 1 cycle, perfect ITLB; 8-wide Fetch within two 16-byte blocks each cycle, potentially over one taken branch TAGE 1+12 components 15k-entry total ( $\simeq$ 32KB) [Seznec and Michaud 2006], 20 cycles min. branch mis. penalty, 2-way 8K-entry BTB, 32-entry RAS 8-wide Decode 8-wide Rename + <b>8-wide Early Execution</b>
Execution	8-wide Dispatch to: 192-entry reorder buffer, 60-entry unified scheduler (IQ), 72-entry load queue, 48-entry store queue 4K-entry SSID/LFST store sets [Chrysos and Emer 1998] memory dependency predictor 256/256 INT/FP—64/128-bit SIMD registers are allocated from the pool of 64-bit FP registers  Each cycle, issue one instruction to 6 ports out of the following 8: Port0: ALU(1c) $\otimes$ Mul(3c) $\otimes$ Div(25c*) $\otimes$ FP Mul(5c) $\otimes$ FP Div(10c*) $\otimes$ 128-bit SIMD FP Mul(5c)Div(25c*) $\otimes$ 128-bit SIMD INT Mul(3c) Port1: ALU(1c) $\otimes$ FP(3c) $\otimes$ 128-bit SIMD ALU (1c INT, 3c FP) Port2: Ld/Str Port3: Ld/Str Port4: ALU(1c) Port6: ALU(1c) Port7: Str  8-wide Writeback 8-wide Validation (only when VP is used) + <b>8-wide Late Execution</b> 8-wide Commit
Caches	L1D 8-way 32KB, 4 cycles, 64 MSHRs, 2 reads and 2 writes/cycle, perfect D-TLB Unified L2 16-way 1MB, 12 cycles, 64 MSHRs, no port constraints L2 Stride prefetcher, degree 8 All caches have 64B lines and LRU replacement
Memory	Single-channel DDR3-1600 (11-11-11), 2 ranks, 8 banks/rank, 8K row-buffer, tREFI 7.8us; min. read lat.: 75 cycles, max. 185 cycles.

\*Not pipelined.

Note: Additional hardware for EOLE is shown in bold.

We consider a relatively aggressive 4GHz, 6-wide issue superscalar<sup>8</sup> baseline with a fetch-to-commit latency of 19 cycles. Since we focus on the out-of-order engine complexity, both in-order front-end and in-order back-end are overdimensioned to treat up to eight  $\mu$ -ops per cycle. We model a deep front-end (15 cycles) coupled to a shallow back-end (3 cycles) to obtain realistic branch/value misprediction penalties.

Table I describes the characteristics of the baseline pipeline that we use in more details. In particular, the out-of-order engine is dimensioned with a unified centralized 60-entry scheduler (also known as an instruction queue (IQ)) and a 192-entry ROB on par with that of Haswell, the latest commercially available Intel microarchitecture. We refer to this baseline as the *Baseline\_6\_60* configuration (6-issue, 60-entry IQ). Functional units are grouped by stacks to mimic the behavior of Haswell's issue ports as best as possible. Note that there are 8 stacks, but since we consider a 6-issue processor, at most 6 stacks will be assigned an instruction in each cycle.

<sup>8</sup>On our benchmark set and with our baseline simulator, an 8-issue machine achieves only marginal speedup over this baseline. Hence, we consider an issue width of 6 to ensure that reducing it will noticeably decrease performance.

In that context, reducing the issue width can be achieved in two fashions. First, it can be reduced by simply removing some issue ports. The associated FUs can then be redistributed to other ports, or be totally discarded from the design. Second, it can be reduced by keeping the same FU stacks but only allowing *new\_issue\_width* instructions to be issued per cycle. An actual implementation will likely do the former; however, since finding the optimal mix of FUs per issue port is beyond the scope of this article, we keep the same FU stacks for all of the different issue widths that we consider in our experiments.

As  $\mu$ -ops are known at Fetch in *gem5*, all widths given in Table I are in  $\mu$ -ops, even for the fetch stage. Independent memory instructions (as predicted by the store sets predictor [Chrysos and Emer 1998]) are allowed to issue out-of-order. Entries in the scheduler are released upon issue, except for load instructions, which release their entry at Writeback.

In the case where VP is used, we add a precommit stage responsible for VT and LE when relevant: the LE/VT stage. This accounts for an additional pipeline cycle (20 cycles) and an increased value misprediction penalty (21 cycles min.). Minimum branch misprediction latency remains unchanged except for mispredicted very high confidence branches when EOLE is used. Note that the value predictor is really trained after Commit, but the value is read from the PRF in the LE/VT stage.

*Shortcomings of the model.* Contrary to modern x86 implementations, *gem5* does not support *move elimination* [Fahs et al. 2005; Jourdan et al. 1998; Petric et al. 2005],  *$\mu$ -op fusion* [Gochman et al. 2003] and does not implement a *stack engine* [Gochman et al. 2003]. It also lacks *macro-op fusion*, *zero-idiom elimination* (e.g., *xor rax, rax*), and a  *$\mu$ -op cache* [Intel 2014].

Although breaking true data dependencies through VP can overlap with some of these features (e.g., zero-idiom elimination and stack engine), most of them are orthogonal; therefore, the inherent increase in ILP brought by VP should not generally be shadowed by said optimizations. In other words, the pipeline that we model is sufficiently resembling a modern superscalar processor to illustrate how VP can push performance higher and how it can help to reduce the aggressiveness of the execution engine.

In addition, there is a limitation with regard to prefetching. Indeed, one could argue that implementing a better prefetcher (we use a simple Stride-based prefetcher in the LLC) could greatly reduce the gains brought by VP through transforming a predicted load that misses in the LLC into a predicted load that hits in the L1.

Although this is true and a better prefetcher may overlap with VP, on our benchmark set we found that on average, only 5% of the loads that miss in the LLC are value predicted. This suggests that even if prefetching could capture all value-predicted loads missing in the LLC, VP would still be able to bring significant performance improvements.

In addition, we point out that correctly predicting a L1 hit would still potentially save four cycles in Haswell but would more importantly allow more freedom regarding speculative scheduling [Kim and Lipasti 2004]. In other words, since load dependents can use the prediction to execute, there is no need to schedule them speculatively, assuming that the load will hit and no bank conflict will take place in the L1 (if the L1 is banked). As a result, scheduling replays due to L1 misses and L1 bank conflicts (hit in L1 but the cache bank is busy) can be avoided if the aforementioned load is value predicted.

## 4.2. Value Predictor Operation

The predictor makes a prediction at fetch time for every eligible  $\mu$ -op (we define *eligible* in the next paragraph). To index the predictor, we XOR the PC of the x86\_64 instruction

with the  $\mu$ -op number inside the `x86_64` instruction. This avoids all  $\mu$ -ops mapping to the same entry for `x86` instructions generating more than one  $\mu$ -op. We assume that the predictor can deliver as many predictions as requested by the fetch stage.

In previous work, a prediction is written into the PRF and replaced by its nonspeculative counterpart when it is computed in the out-of-order engine [Perais and Seznec 2014b]. In parallel, predictions are put in a FIFO queue to be able to validate them—in order—at commit time. In EOLE, we also use a queue for validation. However, instead of directly writing predictions to the PRF, we place predictions in the EE units, which will in turn write the predictions to the PRF at Dispatch. By doing so, we can use predictions as operands in the EE units.

*Eligible  $\mu$ -ops.* Since the predictor is able to produce 64-bit values, all  $\mu$ -ops writing to a 64-bit (or less) general-purpose register (GPR) are predicted, including instructions that convert an FP value to an integer value (e.g., *Convert Single Scalar FP to Signed D/Quadword*, *CVTSS2SI*).

In addition, *gem5-x86* splits 128-bit packed instructions into two 64-bit (32-bit packed or 64-bit scalar)  $\mu$ -ops.<sup>9</sup> As a result, it appears possible to predict FP results (both scalar and vector), as well as packed integer results, and in our first experiments, we consider doing so. However, cracking packed instructions into several  $\mu$ -ops is not representative of a high-performance implementation. Thus, predicting packed 128-bit (or more) registers would in fact involve looking up several predictor entries for the same  $\mu$ -op, or provisioning 128 bits (or more) in each predictor entry, which may not be practical. Consequently, our final experiments do not consider packed instructions as a target for VP.

Nonetheless, scalar FP operations produce a 64-bit (or 32-bit for single-precision FP) value; therefore, it should be possible to predict them. Unfortunately, both encodings (legacy and VEX since AVX) of the `x86` FPU, SSE, require that a scalar operation on a SIMD register copies the upper part of the first source *xmm* register to the destination *xmm* register.<sup>10</sup> For instance, a scalar double-precision FP addition can be expressed by the following algorithm, assuming 256-bit wide *yymm* SIMD registers<sup>11</sup> (as implemented in Haswell [Intel 2013]):

*Legacy:*

```
dest[63:0] = dest[63:0] + src[63:0]
dest[255:64] = dest[255:64] (unmodified)
```

*VEX encoded:*

```
dest[63:0] = src1[63:0] + src2[63:0]
dest[127:64] = src1[127:64]
dest[255:128] = 0
```

Consequently, although the functional unit computes a 64-bit value, the need to merge the physical destination register with one physical source register yields an effective result width of 256 bits (Legacy) and 128 bits (VEX encoded). Hence, except for scalar loads (which zero the upper part of both *xmm* and *yymm*), it is not useful to predict scalar FP results,<sup>12</sup> because the RAW dependency is not broken by the prediction.

<sup>9</sup>Both can be executed in a single cycle in our model, and hence execution throughput is still 128 bits per cycle for packed instructions.

<sup>10</sup>In the legacy encoding, the first source is also the destination register, but the copy of the upper part of the old physical register to the new one is still required.

<sup>11</sup>128-bit *xmm* registers correspond to the lower 128 bits of 256-bit *yymm* registers.

<sup>12</sup>At least for SSE/AVX, which is the current FPU ISA for `x86`, `x87` being deprecated [Intel 2013].



Table II. Layout Summary

Predictor	Entries (#)	Tag	Size (KB)
D-VTAGE [Perais and Seznec 2015]	8,192 (Base) 6 × 1024 (Tagged)	— 12 + <i>rank</i>	131.0 62.6

Note: *rank* is the position of the tagged component and varies from 1 to 6, with 1 being the component using the shortest history length.

Note, however, that this limitation is purely inherent to the ISA and that it would be possible to overcome it at the microarchitectural level, such as by allowing different parts of a single *ymm/xmm* register to be renamed to different physical registers and injecting a *merge*  $\mu$ -op where appropriate.<sup>13</sup> Since this would allow to eliminate many copies of the upper part of *xmm* registers, this might in fact be an optimization that is already implemented. In that event, predicting scalar FP instructions would be possible.

To summarize, whereas we consider predicting all results (i.e., scalar/vector integer and scalar/vector FP) as a first step, to gauge the potential for additional performance and the predictability of FP results, we will ultimately consider predicting values that are written to a GPR only, as predicting scalar FP operations requires microarchitectural support due to ISA limitations, and FP load coverage tends to be very low, on our benchmark set.

*x86 Flags.* In the x86\_64 ISA, some instructions write flags based on their results, whereas some need them to execute (e.g., conditional branches) [Intel 2013]. We assume that flags are computed as the last step of VP, based on the predicted value. In particular, the *Zero Flag*, *Sign Flag*, and *Parity Flag* can easily be inferred from the predicted result. Remaining flags—*Carry Flag*, *Adjust Flag*, and *Overflow Flag*—depend on the operands and cannot be inferred from the predicted result only. We found that always setting the *Overflow Flag* to 0 did not cause many mispredictions and that setting *Carry Flag* if *Sign Flag* was set was a reasonable approximation. The *Adjust Flag*, however, cannot be set to 0 or 1 in the general case. This is a major impediment to the value predictor coverage, since we consider a prediction as incorrect if one of the derived flags—thus the flag register—is wrong. Fortunately, x86\_64 forbids the use of decimal arithmetic instructions. As such, *Adjust Flag* is not used, and we can simply ignore its correctness when checking for a misprediction [Intel 2013].

*Predictor considered in this study.* In this study, we focus on the tightly coupled hybrid of VTAGE and a simple Stride-based predictor, D-VTAGE [Perais and Seznec 2015]. For confidence estimation, we use FPC [Perais and Seznec 2014b]. In particular, we use 3-bit confidence counters whose forward transitions are controlled by the vector  $v = \{1, \frac{1}{16}, \frac{1}{16}, \frac{1}{16}, \frac{1}{16}, \frac{1}{32}, \frac{1}{32}\}$ , as we found it to perform best with D-VTAGE. Those counters allow to push accuracy very high while only requiring 3-bits per counter (the linear feedback shift register used to provide randomness is amortized on the whole predictor, or a large number of entries). This is required to absorb the cost of validating predictions at Commit and squashing to recover.

We consider an 8K-entry base predictor and 6 1K-entry partially tagged components. We do not try to optimize the size of the predictor (i.e., by using partial strides), but it has been shown that 16KB to 32KB of storage are sufficient to get good performance with D-VTAGE [Perais and Seznec 2015].

<sup>13</sup>This is also a way to handle x86 general-purpose partial registers.

Table III. Benchmarks Used for Evaluation

Program	Input	IPC
164.zip (INT)	input.source 60	0.835
168.wupwise (FP)	wupwise.in	1.337
171.swim (FP)	swim.in	2.206
172.mgrid (FP)	mgrid.in	2.356
173.aplu (FP)	aplu.in	1.481
175.vpr (INT)	net.in arch.in place.out dum.out - nodisp -place_only -init_t 5 -exit_t 0.005 -alpha_t 0.9412 -inner_num 2	0.664
177.mesa (FP)	-frames 1000 -meshfile mesa.in - ppmfile mesa.ppm	1.268
179.art (FP)	-scanfile c756hel.in -trainfile1 a10.img -trainfile2 hc.img -stride 2 -startx 110 -starty 200 -endx 160 -endy 240 -objects 10	0.441
183.quake (FP)	inp.in	0.655
186.crafty (INT)	crafty.in	1.551
188.amm (FP)	amm.in	1.212
197.parser (INT)	ref.in 2.1.dict -batch	0.872
255.vortex (INT)	lendian1.raw	1.795
300.twolf (INT)	ref	0.468
400.perlbench (INT)	-I./lib checkspam.pl 2500 5 25 11 150 1 1 1 1	1.370
401.bzip2 (INT)	input.source 280	0.780
403.gcc (INT)	166.i	1.028
416.gamess (FP)	cytosine.2.config	2.196
429.mcf (INT)	inp.in	0.116
433.milc (FP)	su3imp.in	0.499
435.gromacs (FP)	-silent -deffnm gromacs -nice 0	0.792
437.leslie3d (FP)	leslie3d.in	2.139
444.namd (FP)	namd.input	2.347
445.gobmk (INT)	13x13.tst	0.845
450.soplex (FP)	-s1 -e -m45000 pds-50.mps	0.271
453.povray (FP)	SPEC-benchmark-ref.ini	1.519
456.hmm (INT)	nph3.hmm	2.016
458.sjeng (INT)	ref.txt	1.302
459.GemsFDTD (FP)	/	2.146
462.libquantum (INT)	1397 8	0.460
464.h264ref (INT)	foreman_ref_encoder_baseline.cfg	1.127
470.lbm (FP)	reference.dat	0.373
471.omnetpp (INT)	omnetpp.ini	0.309
473.astar (INT)	BigLakes2048.cfg	1.166
482.sphinx3 (FP)	ctlfile . args.an4	0.787
483.xalancbmk (INT)	-v t5.xml xalanc.xml	1.934

Note: Top, CPU2000; bottom, CPU2006; INT, 18; FP, 18; total, 36.

### 4.3. Benchmarks

We use a subset of the the SPEC'00 [Standard Performance Evaluation Corporation 2000] and SPEC'06 [Standard Performance Evaluation Corporation 2006] suites to evaluate our contribution as we focus on single-thread performance. Specifically, we use 18 integer benchmarks and 18 FP programs.<sup>14</sup> Table III summarizes the benchmarks

<sup>14</sup>We do not use the whole suites due to some currently missing system calls or instructions in *gem5-x86*.




Fig. 8. Speedup over *Baseline\_6\_60* brought by VP using D-VTAGE. Instructions eligible for VP are either those writing a GPR register (GPR) or all instructions writing a register (GPR + xmm).

that we use, as well as their input, which are part of the *reference* inputs provided in the SPEC software packages. To get relevant numbers, we identify a region of interest in the benchmark using Simpoint 3.2 [Perelman et al. 2003]. We simulate the resulting slice in two steps: first, we warm up all structures (caches, branch predictor, and value predictor) for 50M instructions, then collect statistics for 100M instructions.

Note that the baseline simulator configuration that we use in this work is significantly different from the configuration used in Perais and Seznec [2014a]: Fetch is less aggressive, the last-level cache is smaller, the first-level data cache is slower, and much less complex functional units and loads ports are implemented. Thus, a more realistic pipeline configuration is simulated, but the instructions per cycle (IPC) in some benchmarks is much lower (e.g., *vpr* has 0.664 vs. 1.326 and *art* has 0.441 vs. 1.211). On the contrary, since we use a bigger LQ (72-entry vs. 48-entry in Perais and Seznec [2014a]) to better fit the Intel Haswell pipeline configuration, and since we use a bigger and better implementation of the store sets memory dependency predictor, IPC is higher in other benchmarks (e.g., *gem5* has 2.196 vs. 1.929 and *parser* has 0.872 vs. 0.544). In almost all cases, the gains come from the better memory dependency predictor, and we actually found that in the version of *gem5* used in Perais and Seznec [2014a], the PC of memory instructions is right shifted by 2 before accessing the table, accommodating 4-byte aligned RISC but not x86's CISC.

## 5. EXPERIMENTAL RESULTS

In our experiments, we first use *Baseline\_6\_60* as the baseline to gauge the impact of adding a value predictor only. Then, in all subsequent experiments, we use *Baseline\_6\_60* augmented with the predictor presented in Table II as our performance baseline. We refer to it as the *Baseline\_VP\_6\_60* configuration. Our objective is to characterize the potential of EOLE at decreasing the complexity of the out-of-order engine. We assume that the EE and LE stages are able to treat any group of up to eight consecutive  $\mu$ -ops every cycle. In Section 6, we will consider trade-offs to enable realistic implementations.

### 5.1. Performance of Value Prediction

Figure 8 illustrates the performance benefit of augmenting the baseline processor with D-VTAGE. The first set of bars consider a predictor that only speculates on the value of GPRs. A few benchmarks present interesting potential (e.g., *applu*, *GemsFDTD*, *h264*, and *astar*), some a more moderate potential (e.g., *wupwise*, *swim*, *mgrid*, *mesa*, *perlbench*, *bzip*, *namd*, and *xalancbmk*), and a few others low potential. No slowdown is observed.

In the second set of bars, the predictor is allowed to speculate on the results of instructions writing to 128-bit *xmm* registers, including integer and FP vector

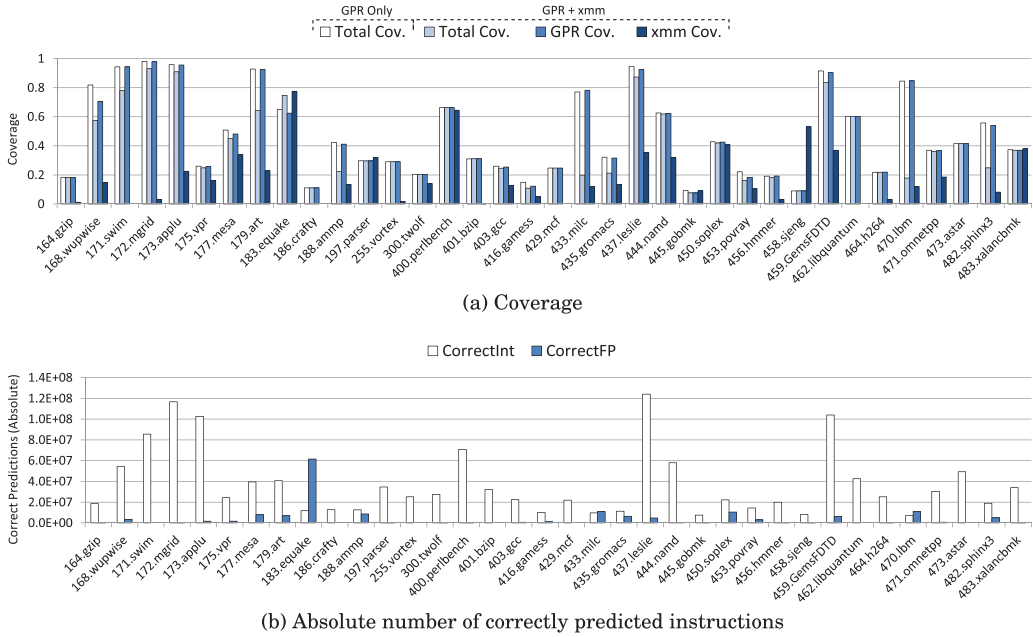


Fig. 9. Relative and absolute coverage of the D-VTAGE predictor.

instructions. As mentioned previously, *gem5-x86* splits a 128-bit instruction into two 64-bit ones; therefore, from the point of view of the predictor, a single packed instruction is in fact two scalar (or 64-bit packed)  $\mu$ -ops. In general, performance is comparable to the previous case, yet in *applu*, *mesa*, *art*, *equake*, *namd*, *soplex*, and *h264*, a slight performance increase can be observed, whereas a slight decrease can be seen in *wupwise*, *gamess*, and *leslie*.

To gain some insight on why we observe such behavior, we consider the overall coverage of the value predictor, as well as the coverage for integer and FP instructions in Figure 9(a). The first observation is that in general, FP coverage<sup>15</sup> is much lower than integer coverage. This is expected since a stride-based prediction scheme cannot naturally predict FP values other than constants. In some cases, however, FP coverage is higher (e.g., *equake*, *gobmk*, *sjeng*, and *xalanbmk*). Yet, since coverage is a relative metric, it is not representative of how many instructions are actually predicted. As a result, Figure 9(b) shows the absolute number of predicted instructions. Except in one benchmark (*equake*), the number of predicted FP instructions is often quite low and, in particular, lower than the number of integer predictions. As a result, in two out of the four previously mentioned cases where FP coverage is higher than INT coverage (*sjeng* and *xalanbmk*), the absolute number of predicted FP instructions is actually very low.

Regardless, benchmarks where performance increases slightly have a similar level of integer coverage, as in the “integer prediction only” case, but they are also able to predict a small (*applu*, *art*, *h264*, *namd*) to moderate (*mesa*, *equake*, *soplex*) amount of FP instructions. On the contrary, when performance decreases, integer coverage is lower than in the “integer prediction only” case, and the small FP coverage does not make up for this reduction. This phenomenon can be explained by the fact that more

<sup>15</sup>FP coverage considers all instructions writing to *xmm*, including integer packed instructions.

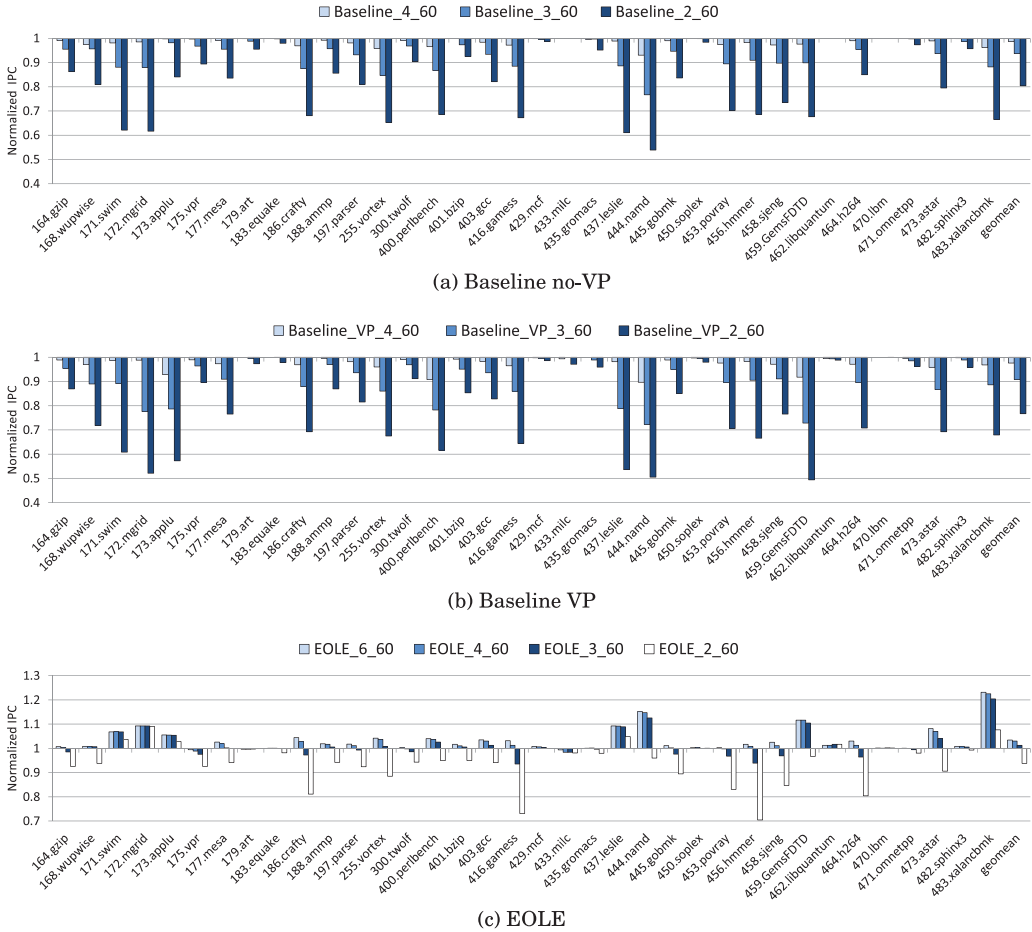


Fig. 10. Impact of a reduced issue width on performance normalized to *Baseline\_6\_60* in (a) and *Baseline\_VP\_6\_60* in (b) and (c).

static instructions have to share the PC-indexed LVT of D-VTAGE, and hence potential decreases because more instructions collide in the LVT.

Nonetheless, it is interesting to note that in *equake*, many FP instructions are actually predictable (the majority being scalar single-precision additions and multiplications but not loads from memory), hinting that although stride-based prediction does not match the FP representation, there are still some FP computations that show enough redundancies to benefit from having a value predictor.

## 5.2. Issue Width Impact on Processor Performance

**Baseline VP model.** We first depict the performance loss implied by reducing the issue width on the baseline processor, *Baseline\_6\_60*, in Figure 10(a). As the issue width is reduced, performance often decreases, but *Baseline\_4\_60* generally performs within 5% of *Baseline\_6\_60*, with a maximum slowdown of 7% in *namd*.

We also experimented with VP (*Baseline\_VP\_6/4/3/2\_60*) in Figure 10(b). It is quite clear that although the average performance of the 4-issue pipeline is close to that of the 6-issue one, up to 10% performance is lost in *applu*, *perlbenc*, *namd*, and *GemFDTD*. Moreover, reducing the issue width beyond 4 is clearly detrimental, as

average performance is 90% that of the 6-issue pipeline for a 3-issue processor and less than 80% for a 2-issue processor (with a maximum slowdown of 50% in *namd* and *GemsFDTD*).

As a result, the reduction in execution engine complexity comes at a noticeable cost in performance when VP is implemented. Without VP, less ILP can be extracted by the processor, and diminishing the issue width is less detrimental, although some benchmarks are still affected as soon as the issue width is decreased to 4.

**EOLE.** The performance of EOLE is illustrated in Figure 10(c), using different out-of-order issue widths.

On top of the 6-issue VP pipeline, EOLE slightly increases performance over the baseline, with a few benchmarks achieving 5% speedup or higher (first bar). The particular case of *namd* is worth to be noted, as with VP only it would have benefited from an 8-issue core by more than 10%. Through EOLE, we actually increase the number of instructions that can be executed each cycle, and hence performance goes up in this benchmark.

For EOLE, reducing the issue width to 4 only decreases performance by a few percent compared with *EOLE\_6\_60*. Furthermore, *EOLE\_4\_60* still performs slightly higher than *Baseline\_VP\_6\_60* in several benchmarks (e.g., *swim*, *mgrid*, *applu*, *leslie*, *namd*, *GemsFDTD*, *astar*, and *xalanmbk*). Slowdowns can be observed in *vpr* and *milc*, but remain modest, at 1.2% and 1.7%, respectively.

Further reduction of the issue width, however, begins to have noticeable impact on performance, with slowdowns of around 5% being observed in *garnet*, *povray*, *hmmmer*, and *h264* on the 3-issue pipeline. As a result, EOLE can be considered as a means to reduce issue width without significantly impacting performance on a processor featuring VP, as long as the execution engine remains wide enough (e.g., 4-issue).

### 5.3. Impact of Instruction Queue Size on Processor Performance

**Baseline VP model.** In Figure 11(a), we illustrate the performance loss inherent to the reduction in the number of IQ entries for the baseline model. Since VP is not present, ILP is often lower, and hence instructions stay longer in the scheduler. This leads to performance decreasing noticeably even when only six IQ entries are removed (e.g., *garnet*, *leslie*, *namd*, *hmmmer*, and *GemsFDTD*), which is not desirable.

Figure 11(b) shows results for the baseline VP pipeline. Since ILP is higher, instructions stay less longer in the IQ, and the performance drop is generally less significant than without VP. However, although the maximum slowdown is less than when the issue width is reduced, almost all benchmarks are slowed down even when only six entries are removed (first bar). On average (gmean), 2% performance is lost, with a maximum of 5.4% in *hmmmer*. Further reducing the IQ size has even more detrimental effects, with an average performance loss of 4% and 8% for a 48-entry and a 42-entry IQ, respectively.

**EOLE.** With EOLE, the performance loss is much less pronounced in general, as shown in Figure 11(c). Yet, even when only six entries are removed from the IQ, a 5% performance degradation is observed in *hmmmer*, which is still problematic, as our goal is to keep at least the same level of performance as the baseline VP model.

In practice, the benefit of EOLE is greatly influenced by the proportion of instructions that are not sent to the out-of-order engine. For instance, *namd* needs a 60-entry IQ in the baseline case, but since it is an application for which many instructions are early or late executed, it can deal with a smaller IQ in EOLE.

On the other hand, *hmmmer*, the application that suffers the most from reducing the IQ size with EOLE, exhibits a relatively low coverage of predicted or early-executed instructions.



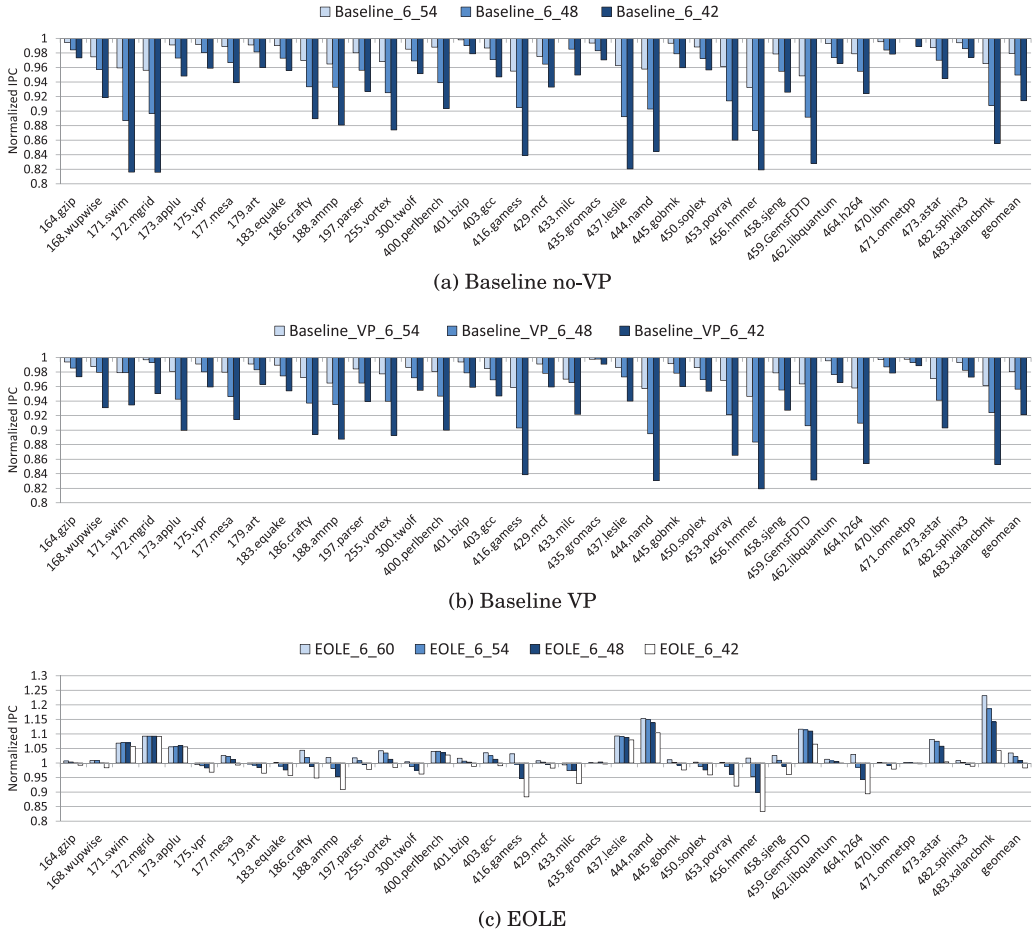


Fig. 11. Impact of a reduced IQ on performance normalized to *Baseline\_6\_60* in (a) and *Baseline\_VP\_6\_60* in (b) and (c).

#### 5.4. Summary

EOLE provides opportunities for either slightly improving the performance over a VP-augmented processor without increasing the complexity of the execution engine or reaching the same level of performance with a significantly reduced execution engine complexity. The said complexity can be reduced via two means: reducing the issue width and reducing the IQ size.

Reducing the issue width reduces the number of broadcast buses used for Wakeup, as well as the complexity of the Select operation. Moreover, the maximum number of ports required on the PRF is reduced, and finally, the number of values co-existing in the bypass network each cycle is also reduced.

Reducing the number of IQ entries, on the contrary, does not reduce the number of broadcast buses, PRF ports, and inflight values inside the bypass network. It only reduces the delay and the power spent in Wakeup by having shorter broadcast buses, as well as the Select delay, as fewer entries have to be scanned [Palacharla et al. 1997]. Regardless, EOLE does not appear as efficient at mitigating the performance

loss inherent to a smaller IQ, and hence reducing the issue width appears as the better choice overall.

In the next section, we provide directions to limit the global hardware complexity and power consumption induced by the EOLE design and the overall integration of VP in a superscalar processor.

## 6. HARDWARE COMPLEXITY

In the previous section, we showed that, provided the processor already implements VP, adopting the EOLE design may allow the use of a reduced-issue execution engine without impairing performance. Yet, extra complexity and power consumption are added in both the EE engine and the LE engine.

In this section, we first describe the potential hardware simplifications on the execution engine enabled by EOLE. Then, we describe the extra hardware cost associated with the EE and LE engines. Finally, we provide directions to mitigate this extra cost. Note, however, that a precise evaluation would require a complete processor design and is beyond the scope of this article.

### 6.1. Shrinking the Out-of-Order Engine

*Out-of-order scheduler.* Our experiments have shown that with EOLE, the out-of-order issue width can be reduced from 6 to 4 without significant performance loss on our benchmark set. This would greatly impact Wakeup, as the complexity of each IQ entry would be lower. Similarly, a narrower issue width mechanically simplifies Select. As such, both steps of the Wakeup & Select critical loop could be made faster and/or less power hungry.

Providing a way to reduce complexity with no impact on performance is also crucial because modern schedulers must support complex features such as *speculative scheduling* and thus *selective replay* to recover from scheduling mispredictions [Kim and Lipasti 2004; Perais et al. 2015].

Last, to our knowledge, most scheduler optimizations proposed in the literature can be added on top of EOLE. This includes the sequential wakeup of Kim and Lipasti [2003] or the tag elimination of Ernst and Austin [2002]. As a result, power consumption and cycle time could be further decreased.

*Functional units and bypass network.* As the number of cycles required to read a register from the PRF increases, the bypass network becomes more crucial. It allows an instruction to “catch” its operands as they are produced and thus execute back-to-back with its producer(s). However, a full bypass network is very expensive, especially as the issue width—hence the number of functional units—increases. Ahuja et al. [1995] showed that partial bypassing could greatly impede performance, even for a simple in-order single-issue pipeline. Consequently, in the context of a wide-issue out-of-order superscalar with a multicycle register read, missing bypass paths may cripple performance even more.

EOLE allows reduction of the issue width in the out-of-order engine. Therefore, it reduces the design complexity of a full bypass by reducing the number of simultaneous writers on the network.

*A limited number of register file ports on the out-of-order engine.* Through reducing the issue width on the out-of-order engine, EOLE mechanically reduces the maximum number of read and write ports required on the PRF for regular out-of-order execution.

### 6.2. Extra Hardware Complexity Associated with Late/Early Execution

*Cost of the late execution block.* The extra hardware complexity associated with LE consists of three main components. First, for validation at commit time, a prediction

queue (FIFO) is required to store predicted results. This component is needed anyway as soon as VP associated with validation at commit time is implemented, since the prediction must be stored until it can be compared against the actual result at Commit.

Second, ALUs are needed for LE. Last, the operands for the late-executed instructions must be read from the PRF. Similarly, the result of VP-eligible instructions must be read from the PRF for validation (predicted instructions only) and predictor training (all VP-eligible instructions).

In the simulations presented in Section 5, we have assumed that up to eight  $\mu$ -ops (i.e., commit width) could be late executed per cycle. This would necessitate eight ALUs and up to 16 read ports on the PRF (including ports required for validation and predictor training).

*Cost of the early execution block.* A single stage of simple ALUs is sufficient to capture most of the potential benefits of EE. The main hardware cost associated with EE is this stage of ALUs and the associated full bypass. Additionally, the predicted and early computed results must be written on the register file.

Therefore, in our case, a complete 8-wide EE stage necessitates eight ALUs, a full 8-to-8 bypass network and eight write ports on the PRF.

*The physical register file.* From the preceding analysis, an EOLE-enhanced core featuring a 4-issue out-of-order engine (*EOLE\_4\_60*) would have to implement a PRF with a total of 12 write ports (respectively, 8 for EE and 4 for regular execution) and 24 read ports (respectively, 8 for regular execution and 16 for LE and VT).

The area cost of a register file is approximately proportional to  $(R + W) * (R + 2W)$ , with  $R$  and  $W$  respectively being the number of read and write ports [Zyuban and Kogge 1998]. In other words, at equal number of registers, the area cost of the EOLE PRF would be four times the initial area cost of the 6-issue baseline without VP (*Baseline\_6\_60*) PRF. Moreover, this would also translate in largely increased power consumption and access time, thus impairing cycle time and/or lengthening the register file access pipeline.

Without any optimization, *Baseline\_VP\_6\_60* would necessitate 14 write ports (respectively, 8 to write predictions and 6 for the out-of-order engine) and 20 read ports (respectively, 8 for VT and 12 for the out-of-order engine) (i.e., slightly less than *EOLE\_4\_60*). In both cases, this overhead might be considered as prohibitive in terms of silicon area, power consumption, and access time.

Fortunately, simple solutions can be devised to reduce the overall cost of the PRF and the global hardware cost of EE/LE without significantly impacting global performance. These solutions apply for EOLE and for a baseline implementation of VP. We describe the said solutions next.

### 6.3. Mitigating the Hardware Cost of Early/Late Execution

*6.3.1. Mitigating the Early Execution Hardware Cost.* Because early-executed instructions are processed in-order and are therefore consecutive, one can use a banked PRF and force the allocation of physical registers for the same dispatch group to different register banks. For instance, considering a 4-bank PRF, out of a group of eight consecutive  $\mu$ -ops, two could be allocated to each bank. In this fashion, a dispatch group of eight consecutive  $\mu$ -ops would at most write two registers in a single bank after EE. Thus, EE would necessitate only two extra write ports on each PRF bank, as illustrated in Figure 12 for an 8-wide Rename/Early Execute, 4-issue out-of-order core. Interestingly, this would add up to the number of write ports required by a baseline 6-issue out-of-order core.

In Figure 13, we illustrate simulation results with a banked PRF. In particular, registers from distinct banks are allocated to consecutive  $\mu$ -ops, and Rename is stalled

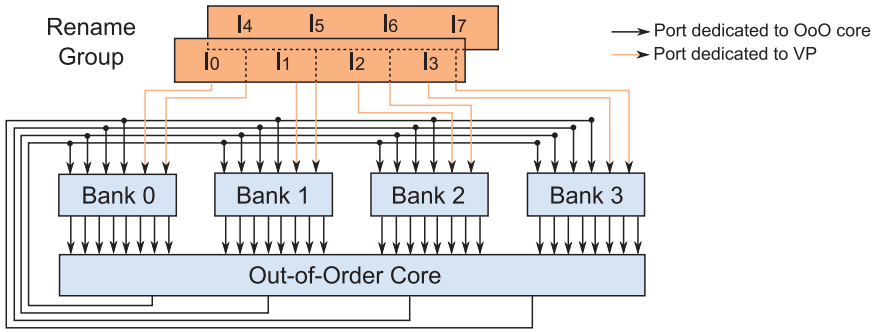


Fig. 12. Organization of a 4-bank PRF supporting 8-wide EE/prediction and 4-wide out-of-order issue. Read ports dedicated to prediction validation and LE are not shown.




Fig. 13. Performance of *EOLE\_4\_60* using a different number of banks in the PRF normalized to *EOLE\_4\_60* with a single bank.

if the current bank does not have any free register. We respectively consider 2 banks of 128 registers, 4 banks of 64 registers, and 8 banks of 32 registers. We observe that the performance loss associated with load unbalancing is quite limited for our benchmark set for the 2- and 4-bank configuration. Therefore, using 4 banks of 64 registers instead of a single bank of 256 registers appears as a reasonable trade-off. However, using 8 banks of 32 registers begins to impact performance, as Rename has to stall more often because there are no free registers in a given bank.

Note that register file banking is also a solution for a practical implementation of a core featuring VP without EOLE, since validation is also done in-order. Therefore, only two read ports are necessary (assuming 4 banks) to validate eight  $\mu$ -ops per cycle.

**6.3.2. Limited Late Execution and Port Sharing.** Not all instructions are predicted or late executable (i.e., predicted and simple ALU or high-confidence branches). Moreover, entire groups of eight  $\mu$ -ops are rarely ready to commit. Therefore, one can limit the number of potentially late-executed instructions and/or predicted IPC. For instance, the maximum commit width can be kept to 8 with the extra constraint of using only six or eight PRF read ports for LE and VT.

Moreover, one can also leverage the register file banking proposed earlier to limit the number of read ports on each individual register file bank at LE/VT. To only validate the prediction for eight  $\mu$ -ops and train the predictor, and assuming a 4-bank PRF, two read ports per bank would be sufficient. However, not all instructions need VT (e.g., branches and stores). Hence, some read ports may be available for LE, although extra read ports might be necessary to ensure smooth LE.




Fig. 14. Performance of *EOLE\_4\_60* (4-bank PRF) when the number of read ports dedicated to LE/VT is limited normalized to *EOLE\_4\_60* (1-bank PRF) with enough ports for full-width LE/VT.

Our experiments showed that limiting the number of read ports on each register file bank dedicated to LE/VT to only four results in a marginal performance loss. Figure 14 illustrates the performance of *EOLE\_4\_60* with a 4-bank PRF and respectively 2, 3, and 4 ports provisioned for the LE/VT stage (per bank). As expected, having only 2 additional read ports per bank is not sufficient. Having 4 additional read ports per bank, however, yields an IPC very similar to that of *EOLE\_4\_60*. Interestingly, adding 4 read ports adds up to a total of 12 read ports per bank (8 for out-of-order execution and 4 for LE/VT)—that is, the same amount of read ports as the baseline 6-issue configuration. Note that provisioning only 3 read ports per bank is also a possibility, as performance is also very close to the ideal *EOLE\_4\_60*.

It should be emphasized that the logic needed to select the group of  $\mu$ -ops to be late executed/validated on each cycle does not require complex control and is not on the critical path of the processor. This could be implemented either by an extra pipeline cycle or speculatively after Dispatch.

**6.3.3. The Overall Complexity of the Register File.** Interestingly, on *EOLE\_4\_60*, the register file banking proposed earlier leads to equivalent performance as a nonconstrained register file. However, the 4-bank file has only 2 extra write ports per bank for EE and prediction and 4 extra read ports for LE/VT. That is a total of 12 read ports (8 for the out-of-order engine and 4 for LE/VT) and 6 write ports (4 for the out-of-order engine and 2 for EE/Prediction), just as the baseline 6-issue configuration without VP.

As a result, if the additional complexity induced on the PRF by VP is noticeable (as issue width must remain 6), *EOLE* allows to virtually nullify this complexity by diminishing the number of ports required by the out-of-order engine. The only remaining difficulty comes from banking the PRF. Nonetheless, according to the previously mentioned area cost formula [Zyuban and Kogge 1998], the total area and power consumption of the PRF of a 4-issue *EOLE* core is similar to that of a baseline 6-issue core *without* VP. Yet, performance is substantially higher since VP is present, as summarized in Figure 15.

It should also be mentioned that the *EOLE* structure naturally leads to a distributed register file organization with one file servicing reads from the out-of-order engine and the other servicing reads from the LE/VT stage. The PRF could be naturally built with a 4-bank, 6 write/8 read ports file (or two copies of a 6 write/4 read ports) and a 4-bank, 6 write/4 read ports one. As a result, the register file in the out-of-order engine would be less likely to become a temperature hotspot than in a conventional design.

**6.3.4. Limiting the Width of Early and Late Execution.** So far, we have considered that *rename\_width* and *commit\_width* ALUs were implemented in the EE stage and the LE stage, respectively. However, for the same reason that the number of read ports dedicated to LE can be limited, the width of EE and LE can be reduced. Indeed, since




Fig. 15. Performance of *Baseline\_VP\_6\_60*, *EOLE\_4\_60* with 16 ports for LE/VT and a single bank, and *EOLE\_4\_60* using 4 ports for LE/VT and having 4 64-register banks normalized to *Baseline\_6\_60*.




Fig. 16. Performance impact of reducing the width of EE/LE. Normalized to *EOLE\_4\_60* with a 4-bank PRF and four read ports dedicated to LE/VT, per PRF bank, and having 8-wide EE/LE stages.

many instructions are not predictable or early/late executable, good enough performance can be attained with only a few additional simple ALUs overall.

Reducing the width of EE is especially interesting, as EE requires a full bypass network to perform best. For instance, reducing the number of ALUs from 8 to 4 reduces the bypass paths from 8-to-8 to 4-to-4. There is a caveat, however, in the fact that predicted results of instructions that are not early executed should also be made available to the next rename group. As a result, bypass is rather 8-to-4 than 4-to-4, as there still might be eight results available for an 8-wide front-end.

Regardless, reducing the width of EE will not cause the pipeline to stall if too many instructions are eligible for EE. Indeed, those instructions will simply have to be executed in the out-of-order engine (unless they are value predicted). On the contrary, reducing the width of LE may stall the pipeline by limiting the number of instructions that can be committed, causing the ROB to become full. For instance, if eight instructions are late executable but only two ALUs are present, then it will take four cycles to commit them, whereas the commit stage could have retired them in a single cycle if eight ALUs had been present.

Note that since we perform prediction validation at LE, we consider that each LE ALU has the comparator required to do so. As a result, by reducing the number of LE ALUs, we also reduce the number of instructions whose prediction can be validated to train the predictor and squash on a misprediction. The maximum commit width of 8 remains attainable if the commit group contains enough instructions that do not produce a register (low-confidence branches that were executed out-of-order, stores, etc.).

Figure 16 depicts the performance impact of decreasing the width of EE/LE. The figure shows IPC normalized to *EOLE\_4\_60* with a 4-bank PRF and four read ports dedicated to LE/VT per PRF bank and having 8-wide EE/LE stages. The first



observation is that for an 8-wide Rename and an 8-wide Commit, 6-wide EE/LE is sufficient, as performance is comparable to having 8-wide EE/LE, except in *namd* (13.3% performance loss). Note, however, that performance in *namd* is still higher than the baseline 6-issue without VP (17.5% speedup).

However, as the EE/LE widths are decreased, performance degrades significantly to attain 70% of the baseline when EE and LE are only 2-wide. Interestingly, using a 2-wide EE stage but a 6-wide LE stage has similar impact as using 6-wide EE/LE stages. This suggests that most of the performance loss is caused by the reduced-width LE stage, as it often limits the actual commit width, whereas the reduced-width EE stage only increases pressure on the execution engine. Consequently, whereas decreasing the EE width from 8 to only 2 does not impact performance significantly, keeping the LE width close to the commit width is necessary to limit pipeline stalls due to the ROB becoming full.

Furthermore, as mentioned previously in Section 3.4.2, a significant amount of value-predicted instructions are loads, which require validation but not LE. Therefore, it is possible that the LE width could be decreased further if LE and prediction VT were to be implemented as two separate pipeline stages and the VT width kept similar to the commit width.

#### 6.4. Late-Executing Predicted Loads

Most processors do not allow more than two loads to issue in the same cycle. The main reason is that adding ports to the data cache to handle more accesses is very expensive. Consequently, banking the data cache is generally preferred, but arbitration increases access latency and the number of bank conflicts will necessarily increase with the number of loads issued each cycle. As a result, it is not desirable to add a data path between the data cache and the LE stage as long as the execution engine already has the ability to issue two loads per cycle.

Nonetheless, by only allowing one load per cycle to be issued by the execution engine, we can move the second load port to the LE stage. If the performance level of the baseline processor can be maintained, doing so may enable further reduction of the out-of-order engine aggressiveness.

Specifically, the load queue (LQ) is implemented to provide correctness in the presence of out-of-order execution of memory operations. It can also be snooped by remote writes to enforce the memory model in some processors (e.g., x86\_64 [Intel 2007]). Since late-executed loads are executed in-order, they will not alias with an older store whose address is unknown, as all older stores have executed, by construction. As a result, it may not be necessary to add late-executed loads to the LQ, depending on the memory model implemented by the processor.

**6.4.1. Late-Executing Loads.** As a first step, and since we focus on a single core, we can ignore the fact that x86\_64 is strongly ordered and consider that loads can safely bypass the LQ while allowing LE not to wait for older loads to return before moving on to younger instructions.<sup>16</sup> This essentially gives us the best of both worlds by potentially allowing reduction of the LQ size while avoiding to stall the pipeline because the ROB became full waiting for a late-executed load to complete, while subsequent instructions could have been late executed in the meantime.

Figure 17 shows IPC normalized to *Baseline\_VP\_6\_60* for different EOLE configurations. The first bar shows performance for *EOLE\_4\_60* with four PRF banks and four read ports dedicated to LE/VT (all of the considered configurations have the same number of PRF banks and read ports). The second bar adds the ability to late execute loads

<sup>16</sup>This can actually lead to a violation of the memory model for x86\_64 [Intel 2007].




Fig. 17. Performance impact of allowing predicted loads to be late executed. IPC is normalized to *Baseline\_VP\_60*. All configurations have four PRF banks and four read ports dedicated to LE/VT.

but can still issue *two* loads out-of-order each cycle, contrary to the third bar, which can only issue *one*. In the two latter cases, validation is able to process up to eight  $\mu$ -ops even if they are not contiguous in the ROB. For instance, when a late-executed load returns, it can be validated in parallel with seven younger  $\mu$ -ops that may be further down the ROB.

Overall, a noticeable performance drop can be observed in several benchmarks when predicted loads are late executed, even if the two loads/cycle throughput is preserved in the out-of-order engine. For the second bar (two loads OoO, one load LE), the reason is that if many loads are predicted, then the maximum load throughput will tend toward one, and since the completion of a predicted load happens later than if it had been executed out-of-order, there is more chance of the ROB becoming full than before. This is the case in *wupwise*, *applu*, and *GemsFDTD*, and this leads to performance dropping even though three loads can theoretically proceed each cycle.

For the third bar (one load OoO, one load LE), if few loads are predicted, then once again, the maximum load throughput will tend toward one, and performance may be reduced (e.g., *gamess* and *h264*).

Consequently, in both low load predictability and high load predictability cases, having a single load port in the out-of-order engine can lead to performance degradation. In the high load predictability case, it would be beneficial to have a steering mechanism deciding if a predicted load is to be late executed or not. This would allow a more efficient use of the two load ports (one in OoO, one in LE). However, in the second case (too few loads are predicted), the only solution to improve the load throughput is to allow nonpredicted loads to be late executed, which is unlikely to be beneficial since it might significantly increase the length of the program critical path.

Moreover, we found that although much fewer instructions enter the scheduler when predicted loads are late executed, neither the scheduler size nor the issue width can be reduced further without impacting performance noticeably. In that context, the only interest of allowing loads to be late executed would be to reduce the LQ size.

**Load queue size.** Figure 18 considers the case where the size of the LQ is varied from the baseline (72-entry) down to 36 entries. Until 48 entries, this has virtually no impact on performance. Even then, although a slowdown is observed in *equake*, 48 entries also provide performance on par with the 72-entry LQ. A 36-entry LQ, however, impedes performance in *equake*, *gamess*, and *h264*. As a result, the LQ should remain large enough to handle cases where few loads are predicted (i.e., the majority of loads must be inserted in the LQ).

In addition, we found that even in the EOLE model that cannot late execute loads, the LQ can be reduced to 60-entry without significant performance loss. This suggests that if not adding late-executed loads to the LQ clearly reduces the pressure put on the structure, the LQ is generally not a bottleneck on our benchmark set.




Fig. 18. Impact of reducing the LQ size on performance when predicted loads are late executed and a single load port is provisioned in the execution engine. IPC is normalized to *Baseline\_VP\_6\_60*.

In any case, in x86\_64, the only actual way to *not* add late-executable loads to the LQ is to enforce that they *complete* in-order in LE. Consequently, to allow some of the latency to be hidden, late-executable loads would have to issue speculatively if an older load were in flight and be replayed if a younger late-executed load were to return before an older one. Due to the numerous events that can lead to such a case (cache misses, cache bank conflicts, bus contention), this may plainly be too impractical for x86\_64. Consequently, we argue that the baseline EOLE model that only allows single-cycle ALU instructions to be early/late executed is a more interesting design point, at least for the window size that is considered.

### 6.5. Summary on the Hardware Complexity of EOLE

Apart from the prediction tables and the update logic, the major hardware overhead associated with implementing VP and validation at commit time comes from the extra read and write ports on the PRF. We have shown earlier that EOLE allows to get rid of this overhead on the PRF as long as enough banks can be implemented.

Specifically, EOLE allows to use a 4-issue out-of-order engine instead of a 6-issue engine. This implies a much smaller instruction scheduler, a much simpler bypass network, and a reduced number of PRF read and write ports in the out-of-order engine. As a result, one can expect many advantages in the design of the out-of-order execution core: significant silicon area savings, significant power savings in the scheduler and the register file, and savings on the access time of the register file. Power consumption savings are crucial, as the scheduler has been shown to consume almost 20% of the power of a modern superscalar core [Ernst and Austin 2002] and is often a temperature hotspot in modern designs. As such, even if global power savings were not to be achieved due to the extra hardware required in EOLE, the power consumption is likely to be more distributed across the core.

On the other hand, EOLE requires some extra but relatively simple hardware for EE/LE. Apart from some control logic, this extra hardware consists of a set of ALUs and a bypass network in the EE stage and a set of ALUs in the LE stage. A full rank of ALUs is actually unlikely to be needed. In fact, from Figure 16, we consider that a rank of two to four ALUs would be sufficient for EE, whereas six to eight ALUs would be required in LE.

In addition, although it would potentially be possible to move one cache port to LE to allow predicted loads to be late executed, we found that this has little practical interest because performance may be reduced due to the uneven number of predicted/nonpredicted loads, as well as the additional pressure put on the ROB by late-executed loads.

Last, implementing EOLE is unlikely to impair cycle time. Indeed, EE requires only one stage of simple ALUs and can be done in parallel with Rename. LE and validation




Fig. 19. Performance of *EOLE\_4\_60*, *OLE\_4\_60*, and *EOE\_4\_60* using four ports for LE/VT and having four 64-register banks normalized to *Baseline\_VP\_6\_60*.

may require more than one additional pipeline stage compared to a conventional superscalar processor, but this should have a fairly small impact, as low-confidence branch resolution is not delayed. In fact, since EOLE simplifies the out-of-order engine, it is possible that the core could actually be clocked higher, yielding even more sequential performance.

Therefore, our claim is that EOLE makes a clear case for implementing VP on wide-issue superscalar processors. Higher performance is enabled thanks to VP (see Figure 15), whereas EOLE enables a much simpler and far less power hungry out-of-order engine. The extra hardware blocks required for EOLE are relatively simple: sets of ALUs in EE and LE stages, and storage tables and update logic for the value predictor itself.

### 6.6. A Note on the Modularity of EOLE: Introducing OLE and EOE

EOLE need not be implemented as a whole. In particular, either EE or LE can be implemented if the performance versus complexity trade-off is deemed worthy. Removing LE can further reduce the number of read ports required on the PRF. Removing EE saves on complexity, as there is no need for an additional bypass network anymore.

Figure 19 shows the respective speedups of *EOLE\_4\_60*, *OLE\_4\_60* (LE only), and *EOE\_4\_60* (EE only) over *Baseline\_VP\_6\_60*. As in the previous paragraph, only four read ports are dedicated to LE/VT, and the PRF is 4-banked (64 registers in each bank). The baseline has a single 256-register bank and enough ports to avoid contention.

We observe that some benchmarks are more sensitive to the absence of LE (e.g., *applu*, *bzip*, *gcc*, *namd*, *hmmr*, and *h264*), whereas one is more sensitive to the absence of EE (*povray*). Nonetheless, removing LE appears as more detrimental in the general case.

However, slowdown over *Baseline\_VP\_6\_60* remains under 5% in all cases. This suggests that when considering an effective implementation of VP using EOLE, an additional degree of freedom exists, as either only EE or LE may be implemented. Indeed, although EE only does not generally allow to increase performance over *Baseline\_VP\_6\_60*, it is still able to maintain the same performance level while the issue width is reduced.

## 7. CONCLUSION AND FUTURE WORK

Single thread performance remains the driving force for the design of high-performance cores. However, hardware complexity and power consumption remain major obstacles to the implementation of new architectural features.

VP is one of such features that has still not been implemented in real-world products due to those obstacles. Fortunately, a recent advance in research on VP partially addressed these issues [Perais and Seznec 2014b]. In particular, it was shown that validation can be performed at commit time without sacrificing performance. This greatly

simplifies design, as the burdens of validation at execution time and *selective replay* for VP in the out-of-order engine are eliminated.

Building on this previous work, we have proposed EOLE, an {Early—Out-of-Order—Late} Execution microarchitecture aiming at further reducing the hardware complexity and the power consumption of a VP-augmented superscalar processor.

With EE, single-cycle instructions whose operands are immediate or predicted are computed in-order in the front-end and do not have to flow through the out-of-order engine. With LE, predicted single-cycle instructions and very high confidence branches are computed in-order in a precommit stage. They also do not flow through the out-of-order engine. As a result, EOLE significantly reduces the number of instructions dispatched to the out-of-order engine.

Considering a 6-wide, 60-entry IQ processor augmented with VP and validation at commit time as the baseline, EOLE allows drastic reduction in the overall complexity and power consumption of both the out-of-order engine and the PRF. EOLE achieves similar or higher performance when using a 4-issue, 60-entry IQ engine, with a few—three on our benchmark set—exceptions being modestly slowed down (1.7% at worst in *milc*).

With EOLE, the overhead over a 6-wide, 60-entry IQ processor (without VP) essentially consists of relatively simple hardware components, the two sets of ALUs in the EE and LE stages, a bypass network and the value predictor tables, and update logic. The need for additional ports on the PRF is also substantially lowered by the reduction in issue width and some PRF optimizations (e.g., banking). Last, the PRF could be distributed into a copy in the out-of-order engine and a copy-only read by the LE/VT stage. Consequently, EOLE results in a much less complex and power-hungry out-of-order engine while generally benefiting from higher performance thanks to VP. Moreover, we hinted that LE and EE can be implemented separately, with LE appearing as more cost effective.

Further studies to evaluate the possible variations of EOLE designs may include the full range of hardware complexity mitigation techniques that were discussed in Sections 6.3 and 6.4 for both EE and LE, and the exploration of other possible sources of LE (e.g., indirect jumps and returns, but also store address computation). One can also explore the interactions between EOLE and previous propositions aiming at reducing the complexity of the out-of-order engine such as the multicluster architecture [Farkas et al. 1997] or register file-oriented optimizations [Wallace and Bagherzadeh 1996]. Finally, future research includes the need to look for even more accurate predictors.

## REFERENCES

- P. S. Ahuja, D. W. Clark, and A. Rogers. 1995. The performance impact of incomplete bypassing in processor pipelines. In *Proceedings of the International Symposium on Microarchitecture*.
- T. M. Austin. 1999. DIVA: A reliable substrate for deep submicron microarchitecture design. In *Proceedings of the International Symposium on Microarchitecture*.
- N. Binkert, B. Beckmann, G. Black, S. K. Reinhardt, A. Saidi, A. Basu, J. Hestness, D. R. Hower, T. Krishna, S. Sardashti, R. Sen, K. Sewell, M. Shoaib, N. Vaish, M. D. Hill, and D. A. Wood. 2011. The gem5 simulator. *SIGARCH Computer Architecture News* 39, 2, 1–7.
- G. Z. Chrysos and J. S. Emer. 1998. Memory dependence prediction using store sets. In *Proceedings of the International Symposium on Computer Architecture*.
- R. J. Eickemeyer and S. Vassiliadis. 1993. A load-instruction unit for pipelined processors. *IBM Journal of Research and Development* 37, 4, 547–564.
- D. Ernst and T. Austin. 2002. Efficient dynamic scheduling through tag elimination. In *Proceedings of the International Symposium on Computer Architecture*.
- B. Fahs, T. Rafacz, S. J. Patel, and S. S. Lumetta. 2005. Continuous optimization. In *Proceedings of the International Symposium on Computer Architecture*.



- K. I. Farkas, P. Chow, N. P. Jouppi, and Z. Vranesic. 1997. The multicluster architecture: Reducing cycle time through partitioning. In *Proceedings of the International Symposium on Microarchitecture*. 11.
- B. Fields, S. Rubin, and R. Bodík. 2001. Focusing processor policies via critical-path prediction. In *Proceedings of the International Symposium on Computer Architecture*.
- F. Gabbay and A. Mendelson. 1998. Using value prediction to increase the power of speculative execution hardware. *ACM Transactions on Computer Systems* 16, 3, 234–270.
- S. Gochman, R. Ronen, I. Anati, A. Berkovits, T. Kurts, A. Naveh, A. Saeed, Z. Sperber, and R. C. Valentine. 2003. The Intel Pentium M processor: Microarchitecture and performance. *Intel Technology Journal* 7, 2, 21–36.
- B. Goeman, H. Vandierendonck, and K. De Bosschere. 2001. Differential FCM: Increasing value prediction accuracy by improving table usage efficiency. In *Proceedings of the International Conference on High-Performance Computer Architecture*.
- Intel. 2007. *Intel 64 Architecture Memory Ordering White Paper*. Retrieved February 21, 2016, from [http://www.cs.cmu.edu/~410/doc/Intel\\_Reordering\\_318147.pdf](http://www.cs.cmu.edu/~410/doc/Intel_Reordering_318147.pdf).
- Intel. 2013. *Intel 64 and IA-32 Architectures Software Developer's Manual*. Available at <http://www.intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-optimization-manual.pdf>.
- Intel. 2014. *Software Optimization Manual*. Retrieved February 21, 2016, from <http://www.fr/content/www/fr/fr/architecture-and-technology/64-ia-32-architectures-optimization-manual.html>.
- S. Jourdan, R. Ronen, M. Bekerman, B. Shomar, and A. Yoaz. 1998. A novel renaming scheme to exploit value temporal locality through physical register reuse and unification. In *Proceedings of the International Symposium on Microarchitecture*.
- R. E. Kessler, E. J. Mclellan, and D. A. Webb. 1998. The Alpha 21264 microprocessor architecture. In *Proceedings of the International Conference on Computer Design*.
- I. Kim and M. H. Lipasti. 2003. Half-price architecture. In *Proceedings of the International Symposium on Computer Architecture*.
- I. Kim and M. H. Lipasti. 2004. Understanding scheduling replay schemes. In *Proceedings of the International Symposium on High-Performance Computer Architecture*.
- M. H. Lipasti and J. P. Shen. 1996. Exceeding the dataflow limit via value prediction. In *Proceedings of the Annual International Symposium on Microarchitecture*.
- M. H. Lipasti, C. B. Wilkerson, and J. P. Shen. 1996. Value locality and load value prediction. In *Proceedings of the International Conference on Architectural Support for Programming Languages and Operating Systems*.
- A. Lukefahr, S. Padmanabha, R. Das, F. M. Sleiman, R. Dreslinski, T. F. Wenisch, and S. Mahlke. 2012. Composite cores: Pushing heterogeneity into a core. In *Proceedings of the International Symposium on Microarchitecture*.
- A. Mendelson and F. Gabbay. 1997. *Speculative Execution Based on Value Prediction*. Technical Report TR1080. Technion-Israel Institute of Technology.
- T. Nakra, R. Gupta, and M. L. Soffa. 1999. Global context-based value prediction. In *Proceedings of the International Symposium on High-Performance Computer Architecture*. 4–12.
- S. Palacharla, N. P. Jouppi, and J. E. Smith. 1997. Complexity-effective superscalar processors. In *Proceedings of the International Symposium on Computer Architecture*.
- A. Perais and A. Seznec. 2014a. EOLE: Paving the way for an effective implementation of value prediction. In *Proceedings of the International Symposium on Computer Architecture*.
- A. Perais and A. Seznec. 2014b. Practical data value speculation for future high-end processors. In *Proceedings of the International Symposium on High-Performance Computer Architecture*.
- A. Perais and A. Seznec. 2015. BeBoP: A cost effective predictor infrastructure for superscalar value prediction. In *Proceedings of the International Symposium on High-Performance Computer Architecture*.
- A. Perais, A. Seznec, P. Michaud, A. Sembrant, and E. Hagersten. 2015. Cost-effective speculative scheduling in high performance processors. In *Proceedings of the International Symposium on Computer Architecture*.
- E. Perelman, G. Hamerly, and B. Calder. 2003. Picking statistically valid and early simulation points. In *Proceedings of the International Conference on Parallel Architectures and Compilation Techniques*.
- V. Petric, T. Sha, and A. Roth. 2005. RENO: A rename-based instruction optimizer. In *Proceedings of the International Symposium on Computer Architecture*.
- B. Rychlik, J. W. Faistl, B. P. Krug, A. Y. Kurland, J. J. Sung, M. N. Velev, and J. P. Shen. 1998. *Efficient and Accurate Value Prediction Using Dynamic Classification*. CM $\mu$ ART-1998-01. Carnegie Mellon University, Pittsburgh, PA.



- Y. Sazeides and J. E. Smith. 1997. The predictability of data values. In *Proceedings of the International Symposium on Microarchitecture*.
- A. Seznec. 2011. Storage free confidence estimation for the TAGE branch predictor. In *Proceedings of the International Symposium on High-Performance Computer Architecture*.
- A. Seznec and P. Michaud. 2006. A case for (partially) TAgged GEometric history length branch prediction. *Journal of Instruction Level Parallelism* 8, 1–23.
- A. Seznec, E. Toullec, and O. Rochecouste. 2002. Register write specialization register read specialization: A path to complexity-effective wide-issue superscalar processors. In *Proceedings of the International Symposium on Microarchitecture*.
- Standard Performance Evaluation Corporation. 2000. SPEC CPU 2000 V1.3. Retrieved February 21, 2016, from <http://www.spec.org/cpu2000/>.
- Standard Performance Evaluation Corporation. 2006. SPEC CPU 2006. Retrieved February 21, 2016, from <http://www.spec.org/cpu2006/>.
- R. Thomas and M. Franklin. 2001. Using dataflow based context for accurate value prediction. In *Proceedings of the International Conference on Parallel Architectures and Compilation Techniques*.
- F. Tseng and Y. N. Patt. 2008. Achieving out-of-order performance with almost in-order complexity. In *Proceedings of the International Symposium on Computer Architecture*.
- E. S. Tune, D. M. Tullsen, and B. Calder. 2002. Quantifying instruction criticality. In *Proceedings of the Conference on Parallel Architectures and Compilation Techniques*.
- S. Wallace and N. Bagherzadeh. 1996. A scalable register file architecture for dynamically scheduled processors. In *Proceedings of the International Conference on Parallel Architectures and Compilation Techniques*.
- K. Wang and M. Franklin. 1997. Highly accurate data value prediction using hybrid predictors. In *Proceedings of the International Symposium on Microarchitecture*.
- H. Zhou, J. Flanagan, and T. M. Conte. 2003. Detecting global stride locality in value streams. In *Proceedings of the International Symposium on Computer Architecture*.
- V. Zyuban and P. Kogge. 1998. The energy complexity of register files. In *Proceedings of the International Symposium on Low Power Electronics and Design*. 305–310.

Received November 2015; accepted December 2015